

Saipriyadarshan Cindula

GenAI Architect · Machine-Learning Expert · 🌐 Hyderabad, India

✉️ priyadarshan.iitk@gmail.com 📞 +91 83413 46818 💎 linkedin/saipriya-cindula



KEY HIGHLIGHTS

- 15+ yrs software, 7+ yrs ML/GenAI
- 100k+ DAU on production LLM stacks
- Multi-agent Temporal/LangGraph expert
- Certified—Databricks GenAI, MS Fabric
- Global delivery across 5 countries
- \$10B+ revenue systems architected
- MCP & A2A protocol specialist
- DeepEval / Arize Phoenix practitioner

CORE SKILLS

Languages

Python JavaScript Java
SQL Bash

GenAI Frameworks

LangGraph LangChain
Temporal DSPy
MCP A2A
AgentFoundry

LLM Providers

OpenAI Anthropic
AWS Bedrock Ollama

RAG & Evaluation

Hybrid RAG DeepEval
Ragas Arize Phoenix
Langfuse W&B

Fine-tuning

LoRA QLoRA PEFT

Claude Ecosystem

Claude API Claude Skills
Claude Code CoWork
Tool Use Prompt Caching
Structured Output

Cloud & Infra

AWS Azure GCP
SageMaker AKS

EXECUTIVE PROFILE

GenAI Architect & ML Engineering Manager with 15+ years in software and 7+ years delivering scalable AI products across BFSI, SaaS, and EdTech. Deep expertise in Advanced RAG, Temporal agentic workflows, multi-cloud (AWS/Azure/GCP), MCP server design, Claude API integration, and xAI evaluation via DeepEval and Arize Phoenix. Track record of translating business KPIs into production architectures serving millions of requests/day.

SELECTED ACHIEVEMENTS & IMPACT

- AgentFoundry / ZenseAI* — Reusable BFSI agent catalogue (digital onboarding, ZenARC regulatory compliance) via MCP, Claude Skills, Claude Code & CoWork; deployed for Standard Bank, NedBank, Investec, Mastercard, EBRD.
- NedBank AI Portfolio Intelligence* — Autonomous treasury agents (LiquidityGuard, CashPredictor, ComplianceWatch); 70% ops reduction, sub-2s decision latency.
- AML/KYC Sanctions Screening* — 5-agent LangGraph pipeline (OFAC, FATF, Basel III); Qdrant hybrid RAG, Ollama LLMs, PostgreSQL audit trail; self-hosted Docker Compose.
- ISO 20022 / SWIFT Migration PoC* — MT101/103 → pain.001/pacs.008 engine for Standard Bank via a 7-layer GenAI Agentic Orchestrator.
- RPG Innovation Festival* — ZenArc AI Compliance Platform across 7 RPG Group subsidiaries (CEAT, KEC, Zensar, RPG Life Sciences, CESC, Spencer's, Harrisons).
- AI Hackathon* — Zenlabs (3rd Prize) — ZenAudit: non-invasive multi-agentic banking compliance with proactive anomaly detection and automated remediation.
- Reliability at Scale* — 99.8–99.99% uptime; 25% cost reduction; \$10B+ revenue systems, zero P0 incidents.

PROFESSIONAL EXPERIENCE

Zensar Technologies

Sr GenAI Architect

Jul 2025–Present
Temporal AI · GenAI Solutions · Scalable Architecture

- Owned BFSI multi-turn bots via Temporal for 1000+ EU & SA users; built MCP server registry (compliance, payments, doc Q&A) with A2A agent-tool binding under Unity Catalog AI governance.
- Architected *AgentFoundry* — reusable foundational BFSI agents (digital onboarding for banks, merchants, capital markets; ZenARC AI-Augmented Regulatory Compliance) composed via MCP, Claude Skills, Claude Code, and CoWork with full-stack Python/React delivery.
- Integrated Claude API prompt caching & structured outputs cutting token spend 40%; deployed Arize Phoenix tracing across A2A hops surfacing hallucination rates and retrieval precision in real time.
- Wired DeepEval BFSI metrics into CI/CD gates; layered NeMo Guardrails for PII and prompt-injection prevention.

NVIDIA Corporation (via Birlasoft)

Senior LLM Engineer

May 2024–Jul 2025
LLM Solutions · Architecture

- Designed Advanced RAG + Hybrid BM25/Dense pipelines for 5k+ users; fine-tuned LoRA/QLoRA/PEFT adapters on OpenAI, Anthropic, and Bedrock; integrated DSPy prompt-tuning cutting manual effort 60%.
- Built DeepEval xAI harness (faithfulness, recall, relevancy) as CI/CD gate; monitored with Arize Phoenix OTel spans; versioned adapters in W&B; ran nightly Ragas benchmarks on domain golden sets.
- Delivered internal knowledge assistant via Claude API multi-doc prompt caching at sub-400 ms p95 latency for 500+ engineers.

Goldman Sachs

Vice President – LLM Engineering & DevOps

Jun 2023–Apr 2024
Full-Stack LLM Development

- Delivered end-to-end LLM apps (Python/FastAPI + React/Next.js) with JWT on Azure AKS; gRPC & WebSockets at <200 ms conversational latency.
- Applied BERT & LIWC for analytics; Pydantic structured outputs for earnings-call extraction; implemented LLM red-teaming guardrail playbook.

GKE EKS

DevOps

Docker K8s

Terraform

GitHub Actions CI/CD

Observability

OpenTelemetry

NeMo Guardrails

Kibana Grafana

Security & Compliance

SOC-2 GDPR

Unity Catalog IAM

EDUCATION

IIT Kanpur

B.Tech Computer Science

2009 · 7.0 / 10

INSEAD

Exec MBA, Technology Leadership

2015 · 92% Distinction

VBCL Hyderabad

LLB Technology & Privacy Law

2025 · First Class

CERTIFICATIONS

- Databricks GenAI Associate
- MS Fabric Lakehouse
- AWS Solutions Architect

LANGUAGES

- English (Fluent)
- Hindi (Native)
- Telugu (Native)

Unify Technologies

Lead LLM Engineer – Technical Delivery

Oct 2022–Jun 2023
Production LLM Systems

- Deployed STT/TTS products (100k+ users, 99.8% uptime) on Azure Kubernetes with full CI/CD; built FAISS embedding recommender, reducing cost 25%.

Alation Inc.

Principal LLM & Data Engineer

Jun 2020–Sep 2022
Scalable LLM Architecture

- Architected distributed metadata platform — 1M+ objects/min on AWS (EKS, RDS, S3); SOC-2 compliant, GDPR-ready with field-level encryption.

Model N Inc.

Principal Software Architect

May 2017–May 2020
Business-Technical Translation

- Built ML pricing models on \$10B+ revenue; multi-region AWS infra (Redshift, VPC, IAM) at 99.99% SLA with blue-green zero-downtime releases.

Next Education India

Senior Software Engineer

May 2015–May 2017
Agile Development

- Led Scrum team; improved outcomes 35% via adaptive testing and NLP question-generator for 50k+ students (TF-IDF & Word2Vec).

InfoEdge (Naukri.com)

Software Engineer

May 2011–May 2015
Production ML Systems

- CTR & salary-prediction at 10k rps; real-time monitoring via Kibana & Grafana with automated drift-triggered model retraining pipelines.

Burning Glass Technologies

ML Engineer

Jul 2009–Apr 2011
Software Development

- CV–Job matching engine (+30% accuracy via gradient-boosted trees) over a canonical 5k-node skill taxonomy powering all downstream analytics.

RESEARCH & ACADEMIC HONOURS

- ★ *Télécom Paris Research Fellow* (2008) — distributed multimedia systems fellowship, top European engineering school.
- ✍️ *ACM Publication* — ACM Multimedia 2008 (pp. 1089–1092).
- ★ *Pratibha Gold Medal & Scholarship* — AP Government award for outstanding academic merit at state level.
- ★ *District Topper & Gold Medalist* — ranked first in engineering at district level; gold medal for academic excellence.
- ★ *NTSE National Scholar* — top 1% in India’s premier science and mathematics talent search programme.
- 🏆 *Hackathon Championships* — multiple first-place finishes in enterprise AI and fintech events.
- *ASOC Ham Radio Licensee* — call sign VU2PDQ; licensed amateur radio operator.
- *BAR Qualified — AIBE 20* — All India Bar Examination qualified; certified to practise law in India.

Portfolio, architecture diagrams & references available on request.